

# Machine Learning Engineer Nanodegree

## Capstone Project

---

Amrit Prasad  
December 19th, 2017

## Definition

---

### Project Overview

On any given day, the stock market moves up or down in an ostensibly random manner. If these random fluctuations can be predicted via a model with more than an even probability, then a trader could utilise said model to profit from the market's moves. For this project, 5 years' worth of historical data relating to S&P 500 Index and VIX were downloaded from Yahoo Finance<sup>[1]</sup>.

### Problem Statement

The aim of this project is to utilise ML techniques, and try to come up with a model that would predict whether S&P 500 will close higher or lower than its open. Since there are two possible outcomes, this can be tackled as a classification problem, and I aim to use a Decision Tree, SVM and Ensemble Learners (Random Forests and Gradient Boosting), and choose the model with the highest cross validation F1 score. A label 1 should imply close above open the next day, while 0 should imply close below open. In case of predicted label 1, go long SPX, while go short in the other case. For unchanged index values, the labelling should be according to the average of the previous five trading days. This prediction can then be used by any amateur trader to profit by taking appropriate positions in SPY ETF/SPX Index Futures/derivatives.

### Metrics

There are a couple of evaluation metrics to consider here-

- 1) F1 score: This is defined as follows-

$$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall}$$

Where,

$$Precision = \frac{True\ Positives}{True\ Positives + False\ Positives}$$

$$Recall = \frac{True\ Positives}{True\ Positives + False\ Negatives}$$

This score is the harmonic mean of how accurately the actual positives have been labelled positive (recall), and how many of the positive labels are actual positives (precision).

- 2) XIRR of the strategy implementation: Once the models have been trained, and the best chosen on the basis of the above, we should try to see the XIRR of running the strategy assuming that the SPX Index can be bought and sold at the index values, which is not too bad an assumption (SPY ETF/Emini Futures do a pretty decent job of capturing the daily moves). The Internal Rate of Return (IRR) is the discount rate that makes the net present value of a series of cashflows zero. It tells you how much a given project/strategy would yield in a given period of time. Greater the IRR, better the strategy. A disadvantage of that is that it's not scaled by the time horizon, which is taken care of by XIRR. XIRR is the annualised IRR of a given set of cashflows. The annualization takes care of the previously discussed issue with IRR.

## Analysis

---

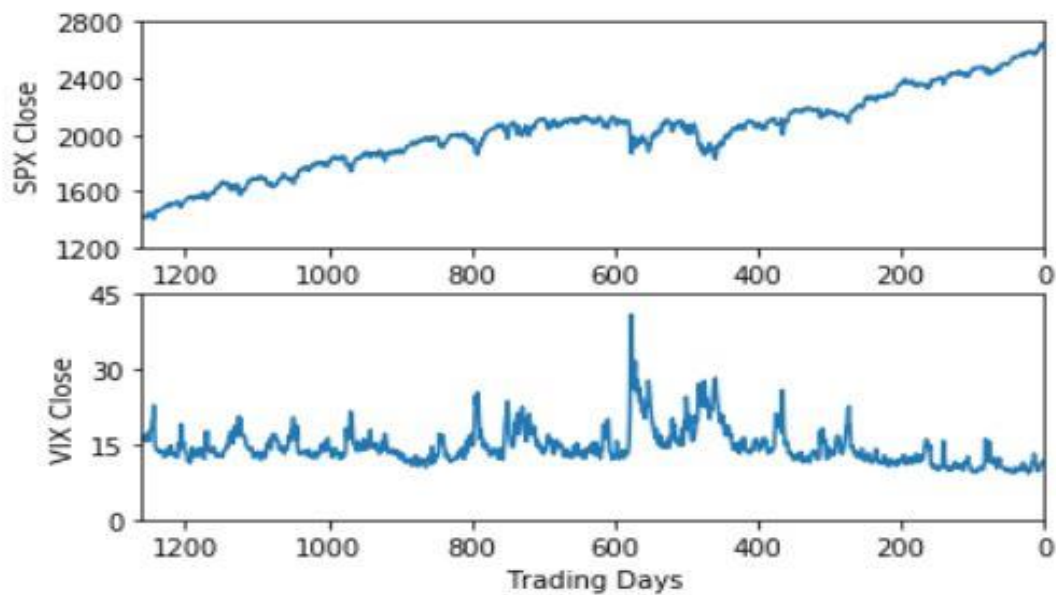
### Data Exploration

For this classification problem, I have tried to limit myself to numerical raw data like open/close price, volume and market expectation of 30-day volatility (VIX). The following raw data fields are going to be used in this problem-

- 1) Date: Date on which the data was recorded
- 2) Open\_SPX: Open value of SPX Index
- 3) High\_SPX: High value of SPX Index
- 4) Low\_SPX: Low value of SPX Index
- 5) Close\_SPX: Close value of SPX Index
- 6) Volume\_SPX: Volume of SPX Index constituents traded
- 7) Open\_VIX: Open value of SPX VIX
- 8) High\_VIX: High value of SPX VIX
- 9) Low\_VIX: Low value of SPX VIX
- 10) Close\_VIX: Close value of SPX VIX

The above dataset was downloaded from Yahoo Finance<sup>[1]</sup> for the last 5 years of trading days from 6<sup>th</sup> Dec 2012 to 6<sup>th</sup> Dec 2017. Each data field has a value corresponding to the Date of trade. If any of the fields is missing/invalid, then it should be filled by the average of the 5 previous trading day values. There are 1260 trading day values corresponding to each of the above-mentioned features in the dataset used.

The following graphs showcase the close of the S&P 500 and VIX Indices-



As can be seen from the graph above, SPX has increased from around 1500 to 2500, while VIX has fluctuated between 5 to 45. Clearly these numerical ranges aren't directly comparable, and new standardised features need to be engineered, that will be the input for training the model. These new features are mentioned in the Project Design section.

As mentioned in the Problem Statement above,  $\text{close} > \text{open} \rightarrow \text{label} = 1$ ,  $\text{close} < \text{open} \rightarrow \text{label} = 0$ , and  $\text{close} = \text{open} \rightarrow \text{label} = \text{average of last 5 labels}$ . If last 5 labels aren't available, then discard the data rows.

## Algorithms and Techniques

Since I have constructed the problem in such a way that the prediction needs to be a binary quantity ( $\text{close} > \text{open}$  or  $\text{close} < \text{open}$ ), and the label is readily calculated, supervised classification techniques would be ideal for attacking this problem. Four such techniques would be used in this project:

- 1) Decision Trees: These are intuitive rule based bifurcators, that split the data on the basis of feature thresholds. Their advantages are their interpretability, speed of training and prediction. Their disadvantages are a tendency to overfit, and high sensitivity to training data.
- 2) SVM: These are extremely versatile classifiers that through the use of kernels can create very accurate models. Their advantages are their accuracy, handling non-linear features, and low susceptibility to outliers since they concentrate only on the data points close to the margin. Their disadvantages are their non-interpretability, slow learning speed, and risk of overfitting in case of low domain knowledge or inability to translate that into the correct choice of kernel.

- 3) Random Forests: These are a kind of Ensemble Learning Method that train multiple weak learners (typically Decision Trees), and then predict on the basis of mode/mean of the individual models' predictions. They tend to reduce the overfitting issue endemic to Decision Trees.
- 4) Extreme Gradient Boosting: Gradient Boosting are another kind of Ensemble Learning Method that creates new learners after weighting the examples which were harder for the previous learner. They are the least prone to overfitting, but can be extremely slow while training, especially if the number of estimators are too high. Extreme Gradient Boosting are a faster implementation of the above, that were chosen because of the very slow training speeds of the Gradient Booster.

## Benchmark Model

The benchmark model to be beaten would be the passive investing strategy in the most popular benchmark index of the stock/index's origin country, which advises buying and forgetting about the investment. This is the benchmark which 80% of fund managers are unable to beat year on year<sup>[3]</sup>. For our case, S&P 500 would be the benchmark index. We can invest in SPY ETFs/long dated forward contracts for simulating the benchmark in real life. For the purposes of this project, I'll be using the SPX Index values as the theoretical benchmark, which should correspond very closely with the real-life version.

## Methodology

---

### Data Preprocessing

As discussed in the Data Exploration section, we need to try and extract useful, standardised information out of the dataset (after it was cleaned as stated in the Data Exploration section), that could be used for predicting whether SPX closes above or below the open. A very important thing to keep in mind while working with this data is to ensure that look-ahead bias is avoided. Hence, if we are trying to predict something on trading day  $x$ , then all data should correspond to trading days  $x-1$  and before. A few useful features that could help us in our prediction for trading day  $x$  are-

- 1) Trailing\_1d\_Return:  $\text{Close}(x-1)/\text{Close}(x-2) - 1$  of SPX Index
- 2) Trailing\_1d\_Max\_Move:  $\text{High}(x-1)/\text{Low}(x-1) - 1$  of SPX Index
- 3) Trailing\_2d\_Return:  $\text{Close}(x-1)/\text{Close}(x-3) - 1$  of SPX Index
- 4) Trailing\_3d\_Return:  $\text{Close}(x-1)/\text{Close}(x-4) - 1$  of SPX Index
- 5) Trailing\_4d\_Return:  $\text{Close}(x-1)/\text{Close}(x-5) - 1$  of SPX Index
- 6) Trailing\_5d\_Return:  $\text{Close}(x-1)/\text{Close}(x-6) - 1$  of SPX Index
- 7) Trailing\_10d\_Return:  $\text{Close}(x-1)/\text{Close}(x-11) - 1$  of SPX Index

- 8) Trailing\_22d\_Return:  $\text{Close}(x-1)/\text{Close}(x-23) - 1$  of SPX Index
- 9) Trailing\_63d\_Return:  $\text{Close}(x-1)/\text{Close}(x-64) - 1$  of SPX Index
- 10) Trailing\_252d\_Return:  $\text{Close}(x-1)/\text{Close}(x-253) - 1$  of SPX Index
- 11) Trailing\_1d\_Return\_VIX:  $\text{Close}(x-1)/\text{Close}(x-2) - 1$  of SPX VIX
- 12) Trailing\_1d\_Max\_Move\_VIX:  $\text{High}(x-1)/\text{Low}(x-1) - 1$  of SPX VIX
- 13) Trailing\_5d\_Return\_VIX:  $\text{Close}(x-1)/\text{Close}(x-6) - 1$  of SPX VIX
- 14) is\_MA5\_above\_MA20: 1 if average of  $\text{Close}(x-1)$  to  $\text{Close}(x-6) >$  average of  $\text{Close}(x-1)$  to  $\text{Close}(x-21)$ ; else 0 of SPX Index
- 15) is\_Trailing\_1d\_Vol\_above\_MA10\_Vol: 1 if  $\text{Vol}(x-1) >$  average of  $\text{Vol}(x-1)$  to  $\text{Vol}(x-11)$ ; else 0 of SPX Index
- 16) Trailing\_1d\_Gap\_Return:  $\text{Open}(x-1)/\text{Close}(x-2) - 1$  of SPX Index

Close(y) above refers to the Close price on trading day y.

Once the above features have been calculated for however many days they can be, the next step involved cleaning up the data by dropping all NaN values, which will creep in if the date ranges are out of scope. Post the removal of NaN, close to 4 years' worth of data was left (the 1-year return was responsible for most of the data rows discarded). After the data cleaning comes the part where we split the data set into training, cross-validation and test sets. The remaining 4 years can then be split into 60%-20%-20%. A key thing to note here is that the data should be split *sequentially* rather than *randomly*, to avoid look-ahead bias.

## Implementation

The first step is to import the evaluation metrics which is the F1 score. The accuracy score is also a good to have.

Post this, the benchmark metrics on the training/cross-validation/test sets were calculated. They were as follows-

Set	F1	Accuracy	XIRR
Training	0.6898	0.5265	5.70%
Cross-Validation	0.6968	0.5347	18.77%
Test	0.6906	0.5274	14.89%

Initially these models were trained on the training dataset to see how they performed while using the default values. Post that they were fine-tuned by changing the hyper parameters. The best tuned model out of them was chosen on the basis of the best F1 score on the cross-validation set. This best model should then be compared against the benchmark on the test dataset using the evaluation metrics of F1 score and XIRR.

## Refinement

This section details the fine tuning that was done to ensure that the models used ended up with the highest possible F1 score on the cross-validation set. The methodology used was fairly straightforward, in that a bunch of critical hyper parameters for each model were recognised, their lists of valid values created, and the models trained repeatedly after passing different values for the parameters.

The following results were observed-

- 1) Two Decision Tree models had better F1 score than the base model. The best one had a maximum depth of 4, and minimum samples per leaf of 9. Its F1 score was 0.6429
- 2) None of the SVMs trained were better than the base model, which calculated that the benchmark was the best fit.
- 3) Eight Random Forests were better than the base model. The best one had a max depth of 1, minimum samples per leaf of 7, split criterion of entropy, 60 number of estimators, and all features used in every constituent tree. Its F1 score was 0.6721
- 4) The Extreme Gradient Booster tuning resulted in the benchmark being calculated as the best fit again.

Best Learner Type	F1	Accuracy	XIRR
Decision Tree	0.6429	0.5545	1.10%
SVM	0.6968	0.5347	18.77%
Random Forest	0.6721	0.5990	62.86%
Gradient Booster	0.6968	0.5347	18.77%

## Results

### Model Evaluation and Validation

None of the models did better on both F1 score and XIRR. In fact, SVM and Extreme Gradient Boosting returned the same values as the benchmark speaking of how hard it is to beat the passive strategy. This tends to verify that the weak form of the efficient market hypothesis holds for S&P 500. The weak form essentially states that all past price/volume movements have been incorporated into an asset's price, and the only way to 'beat the market' is by utilising fundamental analysis. The strong form invalidates the fundamental analysis aspect as well, but verifying that is out of the scope of this project.

Decision Tree and Random Forest were better on the accuracy and XIRR fronts though on the cross-validation set. The better of them, i.e. Random Forest was implemented on the test set as well, after being trained on the combined training + cross-validation sets. It again fails to beat the F1 score of the benchmark, but its XIRR remains exceptional. I would be very sceptical of using this however for trading because it regularly fails to beat the F1 score, indicating the returns are more due to fortuitous trades on the cross-validation/test sets rather than an indication of actual predictive power.

## Reflection

There were two parts of this project that particularly challenged me-

- 1) Getting the cashflows in the correct form of tuples for feeding into the XIRR function that I took from [4]. Joining tuples is a mess. They can be added as `tup_1 + tup_2` or `(tup_1,) + (tup_2,)` or `tup_1, tup_2`. Only one of them gives the correctly accepted form, and googling for this error is a nightmare for a Python neophyte, because of the multiple ways of joining.
- 2) Figuring out what hyper parameters to tune, and get that F1 score to go beyond the benchmark's was the hardest part. Unfortunately, I didn't succeed there. This could be either due to S&P 500 being weak market efficient, or the features chosen by me didn't have enough predictive power.

## Bibliography

---

- [1] <https://finance.yahoo.com/quote/%5EVIX/history?period1=1354818600&period2=1512585000&interval=1d&filter=history&frequency=1d>
- [2] <https://github.com/amritprasad/Capstone/blob/master/Proposal.pdf>
- [3] <http://us.spindices.com/documents/research/research-fleeting-alpha-evidence-from-the-spiva-and-persistence-scorecards.pdf>
- [4] <https://github.com/peliot/XIRR-and-XNPV/blob/master/financial.py>